

Sta 440 Case 3

Srika Gopal, Olivia Fu, Steph Reinke

2025-10-01

1. Background

Emergency medical service (EMS) response times can drastically impact patient outcomes. Vance County, North Carolina is a growing county that is interested in evaluating their current EMS ambulance placements. The county has three major regions, the North, Central, and the South. The population is mostly concentrated in the city of Henderson, located in the Central region. There are currently 2 EMS stations, 1 in the South, and 1 in Central. The southern station has 1 ambulance, and the central station has 3. The north does not currently have a station. Thus Vance County is interested in exploring if moving an ambulance to the North would be beneficial. There are two potential locations for an ambulance to be stationed in the North, the Near North and the Far North. We were given five scenarios for ambulance placement:


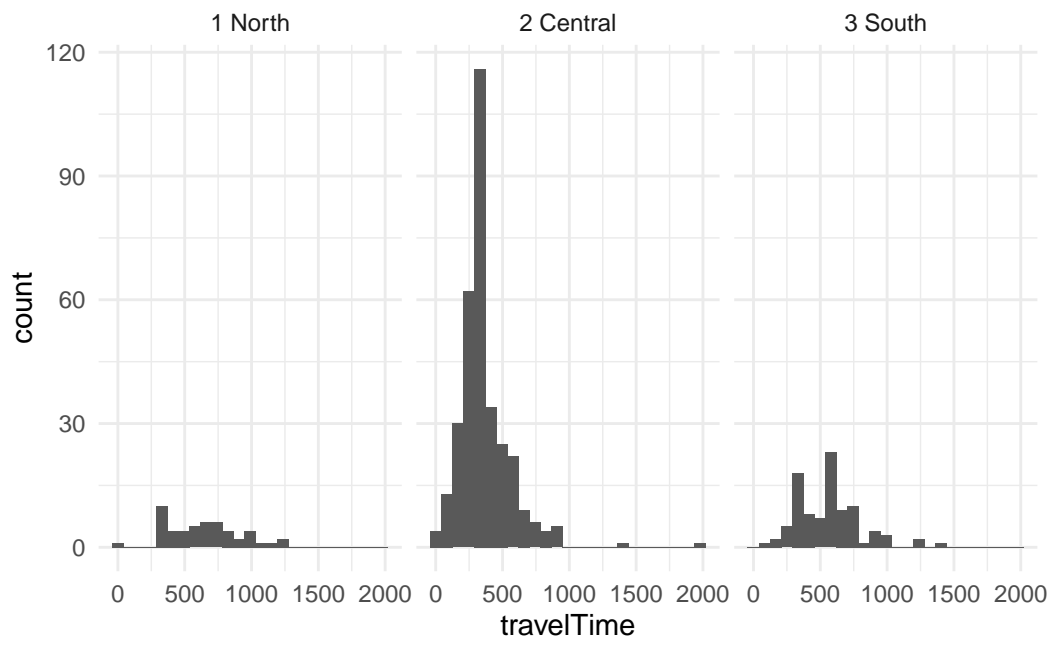
- Scenario 0: 1 in South, 3 in Central (Current Placement)
- Scenario 1: 0 in South, 3 in Central, 1 in Near North
- Scenario 2: 0 in South, 3 in Central, 1 in Far North
- Scenario 3: 1 in South, 2 in Central, 1 in Near North
- Scenario 4: 1 in South, 2 in Central, 1 in Far North


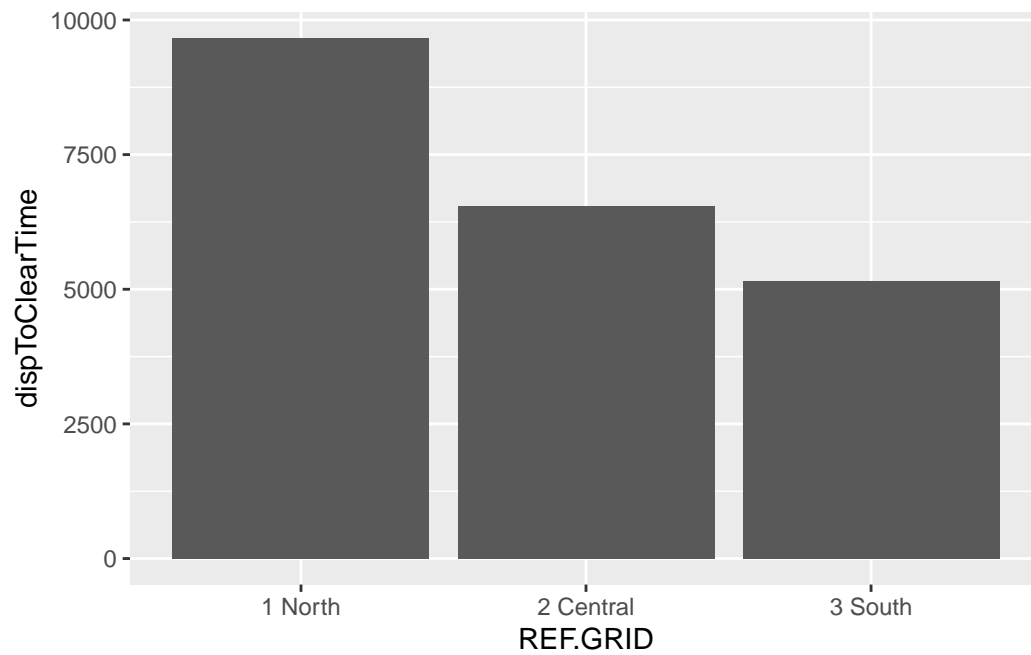
Finding the optimal locations for ambulances is vital because in the medical world EMS arriving seconds earlier can be the difference between life and death.



Research question: Where should the ambulances be stationed to best serve Vance county?

2. Data and Model Fit

To answer this question we examined a dataset with real EMS call data. Select information was withheld or changed, such as dates and addresses for HIPAA compliance. However for the purposes of our analysis those changes are not relevant. The data has information for each call and API pulled travel times from Google Maps. We were provided with four types of estimates from Google Maps (best guess, optimistic, pessimistic, and unadjusted) for each station including the two proposed Northern stations to each call.

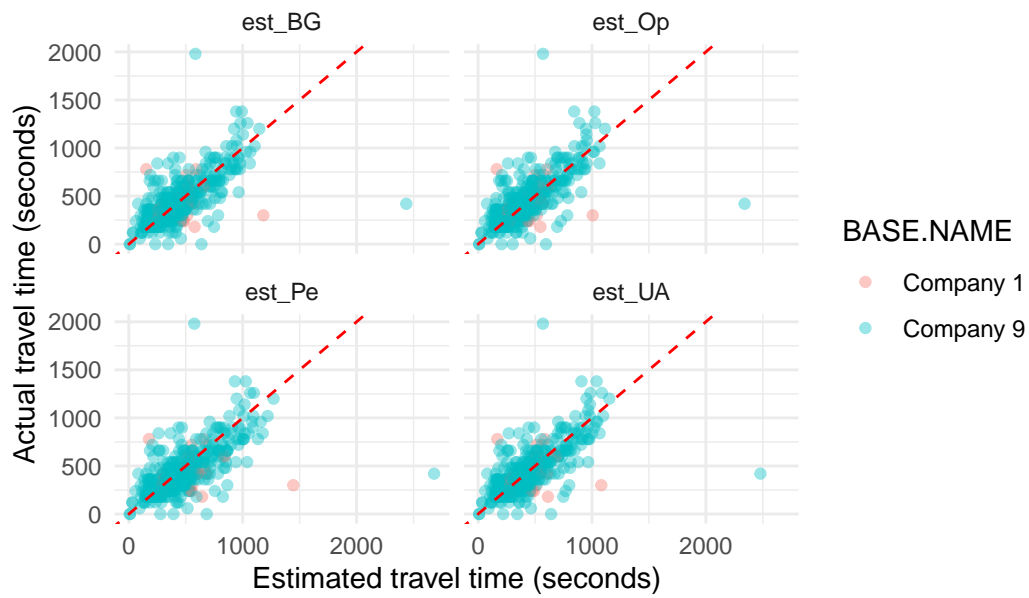




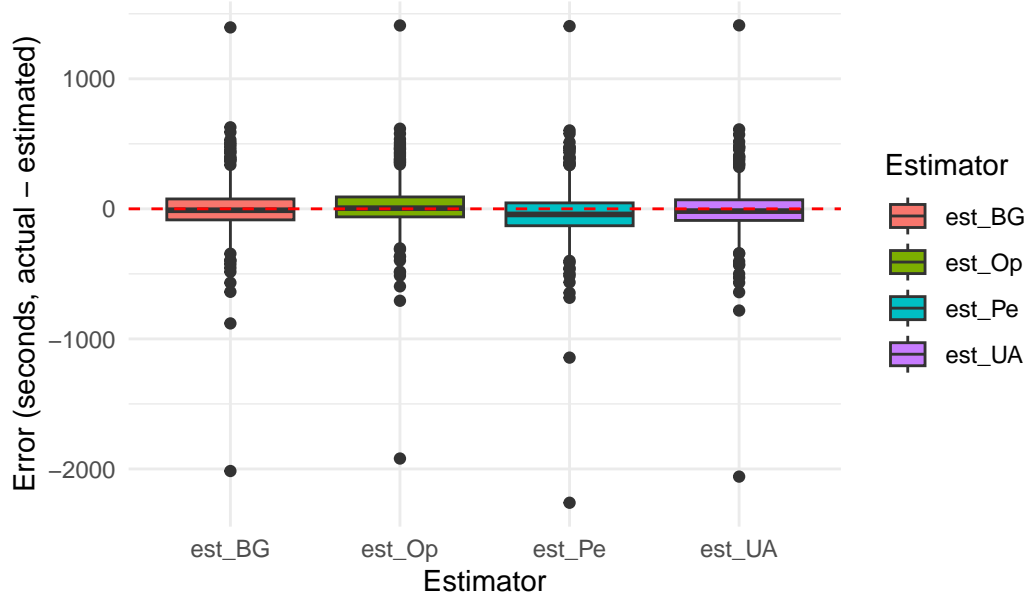


```
# A tibble: 4 x 4
  Estimator MAE RMSE Bias
  <chr>      <dbl> <dbl> <dbl>
1 est_BG    116.  192.  -1.74
2 est_Op    112.  186.   16.3
3 est_Pe    134.  214. -41.7
4 est_UA    116.  192.  -9.14
```

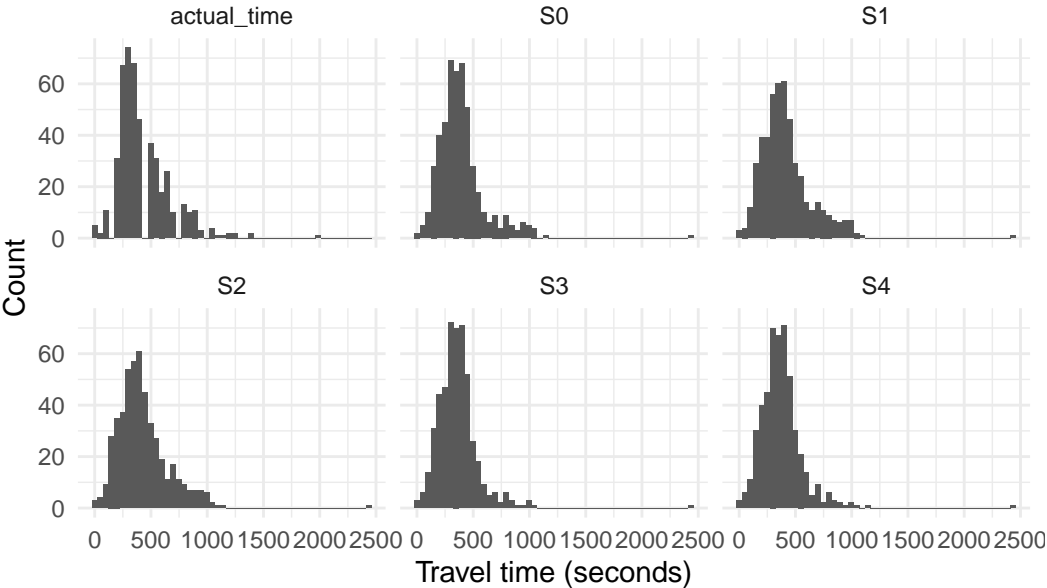
Actual vs. Estimated Travel Times




Error Distribution by Estimator



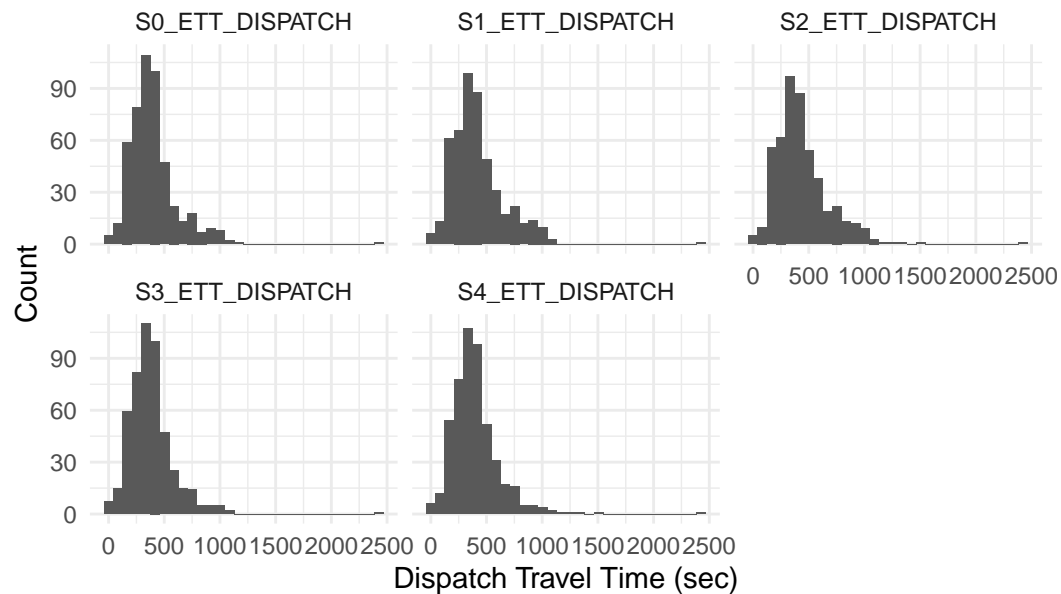
Shortest Travel Time Distributions by Scenario



Average Dispatch ETT per Scenario




Histogram of Dispatch ETT by Scenario



A tibble: 20 x 3


	Scenario	DelayGroup	N
	<chr>	<fct>	<int>
1	S0_ETT_DISPATCH	0-6 min	244
2	S0_ETT_DISPATCH	7-9 min	166
3	S0_ETT_DISPATCH	10-15 min	63
4	S0_ETT_DISPATCH	>15 min	19
5	S1_ETT_DISPATCH	0-6 min	228
6	S1_ETT_DISPATCH	7-9 min	154
7	S1_ETT_DISPATCH	10-15 min	84
8	S1_ETT_DISPATCH	>15 min	26
9	S2_ETT_DISPATCH	0-6 min	212
10	S2_ETT_DISPATCH	7-9 min	158
11	S2_ETT_DISPATCH	10-15 min	95
12	S2_ETT_DISPATCH	>15 min	27
13	S3_ETT_DISPATCH	0-6 min	253
14	S3_ETT_DISPATCH	7-9 min	166
15	S3_ETT_DISPATCH	10-15 min	61
16	S3_ETT_DISPATCH	>15 min	12
17	S4_ETT_DISPATCH	0-6 min	236
18	S4_ETT_DISPATCH	7-9 min	169
19	S4_ETT_DISPATCH	10-15 min	74
20	S4_ETT_DISPATCH	>15 min	13

Counts of Dispatch Delays by Scenario




Scenario S0 S1 S2 S3 S4

Dispatch Delay Counts by Scenario



Scenario S0 S1 S2 S3 S4



Linear mixed model fit by REML ['lmerMod']
 Formula: obsTTmodified ~ (1 | REF.GRID) + DISPATCH.PRIORITY.NAME
 Data: x

REML criterion at convergence: 2560.9

Scaled residuals:

Min	1Q	Median	3Q	Max
-2.9133	-0.5399	-0.2182	0.3473	7.6984

Random effects:

Groups	Name	Variance	Std.Dev.
REF.GRID	(Intercept)	5.617	2.370
	Residual	12.509	3.537

Number of obs: 476, groups: REF.GRID, 3

Fixed effects:

	Estimate	Std. Error	t value
(Intercept)	8.377	1.388	6.037
DISPATCH.PRIORITY.NAMENon Emergency	1.138	0.393	2.896

Correlation of Fixed Effects:

	(Intr)
DISPATCH.PE	-0.059

Generalized least squares fit by REML

Model: obsTTmodified ~ DISPATCH.PRIORITY.NAME
 Data: x

AIC	BIC	logLik
-----	-----	--------

2568.907 2585.552 -1280.454

Correlation Structure: Compound symmetry

Formula: ~1 | region_category

Parameter estimate(s):

Rho

0.3098952

Coefficients:

	Value	Std.Error	t-value	p-value
(Intercept)	8.377225	1.387607	6.037175	0.000
DISPATCH.PRIORITY.NAMENon Emergency	1.137910	0.392978	2.895608	0.004

Correlation:

(Intr)

DISPATCH.PRIORITY.NAMENon Emergency -0.059

Standardized residuals:

Min	Q1	Med	Q3	Max
-2.2348794	-1.0281062	-0.5583537	0.1138833	5.7833057

Residual standard error: 4.257561

Degrees of freedom: 476 total; 474 residual

Call:

```
glm(formula = change_flag ~ region_category + DISPATCH.PRIORITY.NAME,  
     family = binomial, data = model_df)
```

Coefficients:

	Estimate	Std. Error	z value	Pr(> z)
(Intercept)	-1.84130	0.07694	-23.933	< 2e-16 ***
region_categoryNorth	1.86199	0.14458	12.878	< 2e-16 ***
region_categorySouth	3.06905	0.12819	23.942	< 2e-16 ***
DISPATCH.PRIORITY.NAMENon Emergency	-0.46880	0.13877	-3.378	0.000729 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 2965.4 on 2459 degrees of freedom
Residual deviance: 2188.6 on 2456 degrees of freedom
AIC: 2196.6

Number of Fisher Scoring iterations: 4

Linear mixed model fit by REML. t-tests use Satterthwaite's method [
lmerModLmerTest]

Formula: log(ETT_sec) ~ Scenario + region_category + DISPATCH.PRIORITY.NAME +

```
(1 | row_val)
Data: model_df
```

REML criterion at convergence: 2109.5

Scaled residuals:

Min	1Q	Median	3Q	Max
-7.4455	-0.1945	0.0088	0.1643	6.1492

Random effects:

Groups	Name	Variance	Std.Dev.
row_val	(Intercept)	0.23670	0.4865
	Residual	0.07798	0.2793

Number of obs: 2460, groups: row_val, 492

Fixed effects:



	Estimate	Std. Error	df	t value
(Intercept)	5.70728	0.03208	631.72338	177.925
ScenarioS1	0.04018	0.01780	1964.00002	2.257
ScenarioS2	0.08253	0.01780	1964.00002	4.636
ScenarioS3	-0.04064	0.01780	1964.00002	-2.283
ScenarioS4	0.00551	0.01780	1964.00002	0.309
region_categoryNorth	0.37144	0.07479	487.99996	4.967
region_categorySouth	0.44788	0.05744	487.99996	7.798
DISPATCH.PRIORITY.NAMENon Emergency	0.04280	0.05498	487.99996	0.778



	Pr(> t)
(Intercept)	< 2e-16 ***
ScenarioS1	0.0241 *
ScenarioS2	3.79e-06 ***
ScenarioS3	0.0226 *
ScenarioS4	0.7570
region_categoryNorth	9.43e-07 ***
region_categorySouth	3.85e-14 ***
DISPATCH.PRIORITY.NAMENon Emergency	0.4367

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1


Correlation of Fixed Effects:

	(Intr)	ScnrS1	ScnrS2	ScnrS3	ScnrS4	rgn_cN	rgn_cS
ScenarioS1	-0.278						
ScenarioS2	-0.278	0.500					
ScenarioS3	-0.278	0.500	0.500				
ScenarioS4	-0.278	0.500	0.500	0.500			
rgn_ctgryNr	-0.314	0.000	0.000	0.000	0.000		
rgn_ctgrySt	-0.424	0.000	0.000	0.000	0.000	0.173	
DISPATCH.PE	-0.397	0.000	0.000	0.000	0.000	0.015	0.057







Histogram of $\log(\text{model_df\$ETT_sec})$



Residuals



Random Effects



Linear mixed model fit by REML. t-tests use Satterthwaite's method [
lmerModLmerTest]
Formula: ETT_sec ~ Scenario + region_category + DISPATCH.PRIORITY.NAME +
(1 | row_val)
Data: changed_df

REML criterion at convergence: 9613

Scaled residuals:

Min	1Q	Median	3Q	Max
-2.4866	-0.5641	-0.0985	0.5039	3.4132

Random effects:

Groups	Name	Variance	Std.Dev.
row_val	(Intercept)	17449	132.1
	Residual	34793	186.5

Number of obs: 715, groups: row_val, 143

Fixed effects:

	Estimate	Std. Error	df	t value
(Intercept)	426.613	28.211	237.481	15.122
ScenarioS1	79.252	22.059	568.000	3.593
ScenarioS2	134.245	22.059	568.000	6.086
ScenarioS3	-56.573	22.059	568.000	-2.565
ScenarioS4	6.196	22.059	568.000	0.281
region_categoryNorth	-5.336	39.322	139.000	-0.136
region_categorySouth	72.876	29.886	139.000	2.438
DISPATCH.PRIORITY.NAMENon Emergency	46.085	35.611	139.000	1.294

Pr(>|t|)

(Intercept)	< 2e-16 ***
ScenarioS1	0.000356 ***

```

ScenarioS2                2.14e-09 ***
ScenarioS3                0.010585 *
ScenarioS4                0.778913
region_categoryNorth      0.892265
region_categorySouth      0.016012 *
DISPATCH.PRIORITY.NAMENon Emergency 0.197760


```



Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Correlation of Fixed Effects:


	(Intr)	ScnrS1	ScnrS2	ScnrS3	ScnrS4	rgn_cN	rgn_cS
ScenarioS1	-0.391						
ScenarioS2	-0.391	0.500					
ScenarioS3	-0.391	0.500	0.500				
ScenarioS4	-0.391	0.500	0.500	0.500			
rgn_ctgryNr	-0.520	0.000	0.000	0.000	0.000		
rgn_ctgrySt	-0.671	0.000	0.000	0.000	0.000	0.483	
DISPATCH.PE	-0.205	0.000	0.000	0.000	0.000	0.039	-0.013

Residuals vs Fitted

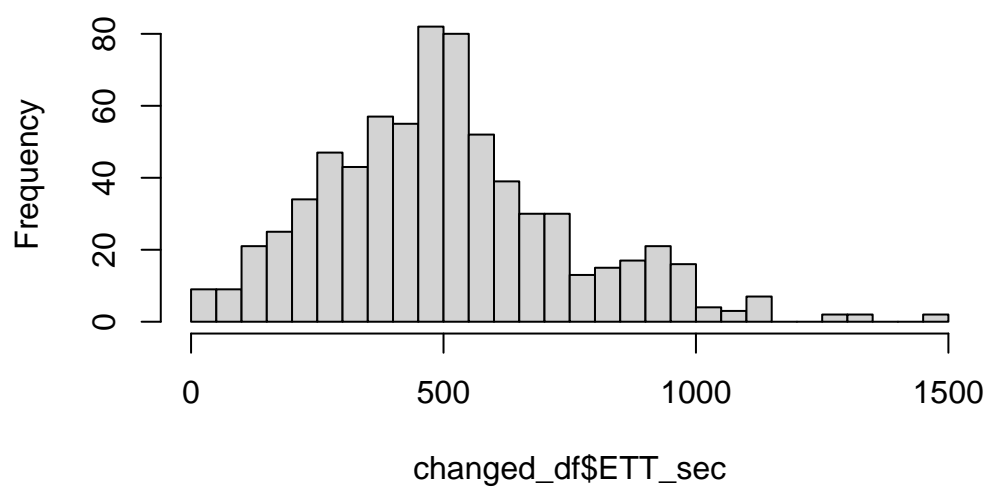







Normal Q-Q Plot




Histogram of changed_df\$ETT_sec







Normal Q-Q Plot



	Model	df	AIC	BIC	logLik	Test	L.Ratio	p-value
gls_model	1	9	33214.66	33266.90	-16598.33			
gls_model_region	2	19	32977.35	33087.57	-16469.67	1 vs 2	257.3084	<.0001

	Model	df	AIC	BIC	logLik	Test	L.Ratio	p-value
gls_model	1	9	33214.66	33266.90	-16598.33			
gls_model_pow	2	18	32954.17	33058.59	-16459.08	1 vs 2	278.4893	<.0001

Type III Analysis of Variance Table with Satterthwaite's method

	Sum Sq	Mean Sq	NumDF	DenDF	F value	Pr(>F)
Scenario	4.2186	1.05465	4	1964	13.5242	7.053e-11 ***
region_category	5.7932	2.89661	2	488	37.1447	9.644e-16 ***
DISPATCH.PRIORITY.NAME	0.0472	0.04725	1	488	0.6059	0.4367

```

Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

      df      AIC
lmm_model 10 2129.54
gls_model  9 33214.66

      df      BIC
lmm_model 10 2187.619
gls_model  9 33266.898

Type III Analysis of Variance Table with Satterthwaite's method

      Sum Sq Mean Sq NumDF DenDF F value    Pr(>F)
Scenario      4.2186  1.05465      4   1964 13.5242 7.053e-11 ***
region_category      5.7932  2.89661      2    488 37.1447 9.644e-16 ***
DISPATCH.PRIORITY.NAME 0.0472  0.04725      1    488  0.6059  0.4367
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

      df      AIC
lmm_model      10 2129.54
gls_model_pow 18 32954.17

      df      BIC
lmm_model      10 2187.619
gls_model_pow 18 33058.592

Type III Analysis of Variance Table with Satterthwaite's method

      Sum Sq Mean Sq NumDF DenDF F value    Pr(>F)
Scenario    3177443  794361      4    568 22.8310 <2e-16 ***
region_category    284976 142488      2    139  4.0953 0.0187 *
DISPATCH.PRIORITY.NAME  58272  58272      1    139  1.6748 0.1978
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

      df      AIC
lmm_model_changed 10 9632.967
gls_model          9 33214.656

      df      BIC
lmm_model_changed 10 9678.69
gls_model          9 33266.90

Type III Analysis of Variance Table with Satterthwaite's method


      Sum Sq Mean Sq NumDF DenDF F value    Pr(>F)
Scenario    3177443  794361      4    568 22.8310 <2e-16 ***
region_category    284976 142488      2    139  4.0953 0.0187 *
DISPATCH.PRIORITY.NAME  58272  58272      1    139  1.6748 0.1978
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

      df      AIC
lmm_model_changed 10 9632.967
gls_model_pow     18 32954.167


```

	df	BIC
lmm_model_changed	10	9678.69
gls_model_pow	18	33058.59

Residuals



Random Effects



	Model	df	AIC	BIC	logLik	Test	L.Ratio	p-value
lmm_no_weights	1	10	2129.540	2187.586	-1054.770			
lmm_model_ext	2	12	1170.544	1240.200	-573.272	1 vs 2	962.9957	<.0001

4. Assumptions and Shortcomings

Normality of residuals was poorly met Large deviances at tails Non-constant variance across fitted values External factors not taken into consideration Traffic / Rush hours/ Road closures Time of year EMT staffing patterns ETT inputs: based on single best-guess estimates (e.g., from Google), lacking nuance or uncertainty Variance modeling: only accounted for regional differences, while other factors may also influence variability Emergency vs non-emergency Simplicity in Dispatch

Rule Better for analysis Doesn't occur in real life

5. Conclusions

Scenario 3 results in the fastest overall response times Marginal mean travel times (ETTs) are reduced by $\sim 16\%$ compared with the baseline (Scenario 0)

$\exp(-0.173) = 0.84 \rightarrow$ travel times in Scenario 3 are about 84% of baseline levels